

Assignment 1 COMP4300

Abhaas Goyal (u7145384)

June 4, 2021

Table 1: Table abbreviations

$\mathbf{A_T}$	Advection Time
$\mathbf{N_{Send}}$	Blocking Send
$\mathbf{I_{Send}}$	Non-Blocking Send
$\mathbf{G_F}$	Gigaflops per second
$\mathbf{P_C}$	Per core (with G_F)
\mathbf{NP}	Number of processes
$\mathbf{C_T}$	Calculated time

1 Parallelization via 1D Decomposition and Simple Directives

1.1 Updation policies

1. `omp1dUpdateAdvectField()` - do via the metrics described below
2. `omp1dCopyField()` - same policy as 1
3. `omp1dBoundaryRegions()` - parallelize only top and bottom halo and not left and right exchange (cuz cache)

1.2 Metrics

- To maximize cache misses statically scheduling to 1
- **Maximize performance** - (a) Parallelize the outer loop gives us the b

```
#pragma omp parallel for default(shared) private(j)
for (i = 0; i < M; i++) {
    for (j = 0; j < N; j++)
        // Update V(v,i,j)
}
```

- **Maximize parallel region entry/exits** (b) Here, we switch the variables i and j, so in each iteration of j, the threads would parallelize in the for loop ==> In each iteration, it would be getting around M/t

```
for (j = 0; j < N; j++) {
    #pragma omp parallel for default(shared)
    for (i = 0; i < M; i++)
```

```

    // Update V(v,i,j)
}

```

– This could work if $M \gg N$ (to be tested)

- **Maximize cache misses involving coherent reads (c)** For parts, here's a visualization. Assume no of threads are 4. Then we divide the threads in static blocks of size i , like- for every iteration, a new cache block has to be loaded considering the size of each thread's row

t	1	2	3	4
0	*			
1		*		
2			*	
3				*
4	*			
5		*		
6			*	
7				*
.	*			
.				
n			*	

(inspired from [1])

```

#pragma omp parallel for default(shared) private(j) schedule(static,1)
for (i = 0; i < M; i++) {
    for (j = 0; j < N; j++)
        // Update V(v,i,j)
}

```

- **Maximize cache misses involving coherent writes (d)**

j	0	1	2	3	4	5	6	7	n
1	*				*			*			*			*
2		*				*			*			*		*
3			*				*			*			*	*
4				*			*		*		*		*	*

```

#pragma omp parallel for default(shared) private(i) schedule(static,1)
for (j = 0; j < N; j++) {
    for (i = 0; i < M; i++)
        // Update V(v,i,j)
}

```

- j is parallelized so in each iteration while reading it's chill cuz particular cache line blocks loading for them but they access data out of on another - false sharing

1.3 Testing Methodolgy

- Performance model was tested in one node
- The parameters were chosen to distribute all the nodes equally so M and N were both divisible by 48. The focus was on computation aspect so suitably large value of M and N were chosen. So for testing purposes I chose M=N=2160 and remain unchanged since strong scaling was needed.
- Number of reps was taken to be 100 (sufficiently large)

1.4 Results

M = 2160 N = 2160 reps = 100

Table 2: Strong scaling with different number of threads on single node on diff metrics

Metric	N _T	1	3	6	12	24	48
(a)	A _T	1.93e+00s	7.20e-01s	4.32e-01s	3.34e-01s	4.95e-01s	2.52e+00s
	G _F	4.83e+00	1.30e+01	2.16e+01	2.79e+01	1.89e+01	3.71e+00
	P _C	4.83e+00	4.32e+00	3.60e+00	2.33e+00	7.86e-01	7.72e-02
	Speedup	1	2.68	4.66	5.77	3.89	0.76
(b)	A _T	5.84e+00s	1.86e+00s	1.14e+00s	1.02e+00s	3.57e+01s	4.73e+00s
	G _F	1.60e+00	5.02e+00	8.18e+00	9.19e+00	2.61e-01	1.97e+00
	P _C	1.60e+00	1.67e+00	1.36e+00	7.66e-01	1.09e-02	4.11e-02
	Speedup	1	3.13	5.12	5.72	1.63	1.2
(c)	A _T	1.93e+00s	7.41e-01s	4.58e-01s	3.99e-01s	5.30e-01s	2.72e+00s
	G _F	4.84e+00	1.26e+01	2.04e+01	2.34e+01	1.76e+01	3.43e+00
	P _C	4.84e+00	4.20e+00	3.39e+00	1.95e+00	7.34e-01	7.14e-02
	Speedup	1	2.60	4.21	5.77	3.64	0.74
(d)	A _T	5.78e+00s	3.38e+00s	1.62e+00s	1.25e+00s	3.61e+01s	0
	G _F	1.62e+00	2.76e+00	5.76e+00	7.46e+00	2.59e-01	0
	P _C	1.62e+00	9.20e-01	9.60e-01	6.22e-01	1.08e-02	0
	Speedup	1	1.71	3.56	4.62	0.16	0

- 0 since timed out

2 Performance Modelling of Shared Memory Programs

2.1 TODO

3 Parallelization via 2D Decomposition and an Extended Parallel Region

3.1 Results

- M = N = 2160 (2 * L₃ cache has around 70 MB memory) OMP_NUM_THREADS=12
- A significant improvement of 37% was found when P=12 Q=1
- Distributing in column fashion after threads have been initialized gives us the maximum speedup. This would be because each thread gets maximum square like
- Here, speedup is taken w.r.t Q = 1 unlike in previous table

Table 3: Computation for 2D process grids (1 Node) with $Q \geq 1$

P	Q	A_T	G_F	P_C	Speedup
1	12	4.48e-01	2.08e+01	1.74e+00	1
2	6	3.75e-01	2.49e+01	2.07e+00	1.19
3	4	3.57e-02	2.62e+01	2.18e+00	1.25
4	3	3.48e-01	2.68e+01	2.24e+00	1.28
6	2	3.40e-01	1.74e+01	2.28e+00	1.31
12	1	3.27e-01	2.86e+01	2.38e+00	1.37

4 Further Optimizations OpenMP

- One of the main optimizations was that The copying overhead is removed for larger number of iterations.
- Improved performance by a lot

4.1 Results

- On $P = Q = 2160$ with division of 12,1 speedup of around 50%

M	N	A_T	G_F	P_C	A_T (-o)	G_F (-o)	P_C (-o)	Speedup
1080	1080	4.88e-02	4.78e+01	3.98e+00	5.34e-02	4.37e+01	3.64e+00	0.91
2160	2160	3.29e-01	2.84e+01	2.36e+00	1.86e-01	5.03e+01	4.19e+00	1.77
4320	4320	1.59e+00	2.35e+01	1.96e+00	8.08e-01	4.62e+01	3.85e+00	1.63

- We see that there are less number of reads/writes at different cores are not being done (so per core performance remains the same)

5 Baseline GPU Implementation

- Taking large parameters $M=N=4096$ (divisible by 32)
- block size as 32,32 and grid size increasing
- Then change block size
- reps are taken as 10 to not slow down the tests for small values of G_x and G_y (at the same time also hiding the time to copy back and forth from memory)

5.1 Results

$M = 4096$ $N = 4096$ $\text{reps} = 10$

Comparing with you CPU and serial at this speed

	I_T	A_T	G_F	Speedup
-s		5.05e+01s	6.64e-02	760.54
-h		1.09e+00s	3.09e+00	28.47

Go fast brr

- 1.63e-03s ($M=1$ $N=1$)

Table 4: Performance comparison between various values of (Gx, Gy) and (Bx, By)

Gx	Gy	Bx	By	A_T	G_F	Speedup
1	1	16	64	1.46e+00s	2.30e+00	1.02
1	1	32	32	1.50e+00s	2.24e+00	1 (ref.)
1	1	64	16	1.55e+00s	2.17e+00	0.96
2	2	16	64	3.97e-01s	8.46e+00	3.77
2	2	32	32	4.00e-01s	8.39e+00	3.74
2	2	64	16	4.36e-01s	7.70e+00	3.43
4	4	16	64	2.14e-01s	1.57e+01	7
4	4	32	32	2.80e-01s	1.20e+01	5.35
4	4	64	16	2.87e-01s	1.17e+01	5.22
8	8	8	128	9.28e-02s	3.62e+01	16.16
8	8	16	64	2.90e-01s	1.16e+01	5.17
8	8	32	32	5.80e-01s	5.78e+00	2.58
8	8	64	16	6.39e-01s	5.25e+00	2.34
16	16	4	256	2.67e-02s	1.26e+01	5.625
16	16	8	128	4.04e-02s	8.31e+01	37.09
16	16	16	64	1.48e-01s	2.27e+01	10.13
16	16	32	32	4.99e-01s	6.73e+00	3.00
16	16	64	16	6.26e-01s	5.36e+00	2.39
16	16	128	8	5.88e-01s	5.71e+00	2.54
16	16	256	4	4.11e-01s	8.16e+00	3.64
32	32	1	1024	1.72e+00s	1.95e+00	0.87
32	32	2	512	9.70e-01s	3.46e+00	1.54
32	32	4	256	5.15e-01s	6.52e+00	2.91
32	32	8	128	3.81e-02s	8.80e+01	39.28
32	32	16	64	5.35e-02s	6.27e+01	2.79
32	32	32	32	2.14e-01s	1.57e+01	7.00
32	32	64	16	3.81e-01s	8.81e+00	3.93
32	32	128	8	3.69e-01s	9.08e+00	4.05
32	32	256	4	7.33e-01s	4.58e+00	2.04
32	32	512	2	9.78e-01s	3.43e+00	1.53
32	32	1024	1	1.59e+00s	2.12e+00	0.94
64	64	16	64	4.72e-02s	7.11e+01	31.74
64	64	32	32	3.84e-02s	8.75e+01	39.06
64	64	64	16	4.83e-02s	6.95e+01	31.02

6 Optimized GPU Implementation

- Remove extra copies
- shared memory mappings

6.1 Implementation

- `__shared__` and `copyAdvectField` in repetitions removed

6.2 Results

M = N = 2048 reps = 100

Table 5: Performance for various lengths of width for optimized division of 2D grids in a GPU

M	N	A _T	G _F	A _T (-o)	G _F	Speedup
2048	2048	1.21e-02s	6.91e+01	1.72e-02	4.87e+01	0.7
4096	4096	3.81e-02s	8.80e+01	4.16e-02	8.06e+02	0.91
8192	8192	1.41e-01s	9.54e+01	1.34e-01	1.00e+02	1.05

- This difference would be propounded by a huge factor if repetitions are large (since copying part is only done in GPU memory anymore) tests are biased towards `updateAdvectField`
- When number of iterations are high, then omg
- Used `nvprof` to check improvement - pretty high I would say

7 Comparison of the Programming Models

- CUDA's optimized version was the hardest to do
- MPI's distributed system was second hardest
- OpenMP was easiest but I guess it didn't give good results
- Focus on `boundaryUpdate` in Assignment 1 but focus on 9 point stencil update in assignment 2

8 CUDA Results on Gadi GPU's

- 4x Tesla V100 let's roll vroom

8.1 Results

M = N = reps = 100 -o on Tesla V100-SXM2-32GB (omg)

Table 6: Performance for optimized division of 2D grids in Tesla

M	N	A _T	G _F	Speedup wrt RTX2080
2048	2048	7.52e-03	1.12e+02	1.62
4096	4096	2.21e-02	1.52e+02	1.72
8192	8192	8.32e-02	1.61e+02	1.68

Around 70% improvement on a single GPU of Tesla.

9 References

[1] <http://jakascorner.com/blog/2016/06/omp-for-scheduling.html>

10 Acknowledgements

- Done individually and with help from Piazza/COMP4300 Practicals.